

Personalised Policy Communication Engine

Kepler Insurance . Case B

AI use case to improve retention through transparent, personalised renewal communications across 3 European markets.

Ground the room: this is a GenAI play to retain premium in a 53%-shopping-rate market while regulators prosecute opaque pricing.

"My renewal shows a higher price every year with no explanation. I don't know why it went up -- so I assume I'm being taken advantage of."

Derived from FCA verbatims: Alan Ford (136% Saga hike) and J.D. Power 2026 finding that only 58% of customers understand their policy, down 4 points from 2025.

Why this matters now: LexisNexis Q1 2026 reports 47.3% annual shop rate. Customers who don't understand their premium are actively shopping. FCA Consumer Duty requires fair-value evidence per renewal. The status quo fails on both retention and compliance.

3. Who Experiences This -- Three User Groups

Renewal Comms Team

- Owns churn and renewal NPS. Currently batch-prints opaque templated letters.
- Cannot personalise at scale -- each market has unique regulatory templates.

Compliance / FCA Consumer Duty

- Needs fair-value evidence per notice. Current manual sampling covers 0.1% of renewals.
- If comms produce opaque letters, compliance can't defend them in regulator review.

Contact Centre

- 23% of inbound call volume is "why did my premium go up?"
- Transparent notices deflected ~40% of these calls in comparable UK motor pilots (2024-2025).

This isn't just a marketing problem -- it's a compliance problem and an ops-cost problem wearing a marketing hat.

Target: Reduce renewal churn 3pp (12% -> 9%) and improve price satisfaction by +30 points within 6 months across first two markets.

Measured via follow-up NPS 30 days post-renewal. Baseline: J.D. Power 2026 -- customers who understand their premium change score +141 points higher on price satisfaction.

W Patch #1: 85% of projected returns depend on Assumption 1 (A/B test). If the A/B shows no retention effect, this value is EUR0.

The J.D. Power correlation is observational, not causal. Satisfaction -> retention conversion is not 1:1. The 3pp target assumes a linear relationship not demonstrated in any published controlled study. We will measure it in the pilot.

A GenAI system integrated into Kepler's renewal notice pipeline generates a personalised, plain-language explanation for every premium change -- deployed in two phases.

Phase 1 -- Decomposition layer (prerequisite)

- Instrument the rating engine's batch pipeline to log every input variable consumed at score time for each policy.
- For legacy PAS books that store only the final premium: fallback template with aggregated market-level attribution.
- Estimated data engineering effort: 2-3 months per market (currently unfunded in base budget).

Phase 2 -- Generation & guardrails

- Claude-class LLM synthesises premium delta + external index feeds + customer profile into a structured renewal letter.
- Compliance rule engine validates: no prohibited language, all numerical claims tied to source data, fair-value narrative internally consistent with pricing decision.

Key design constraint: The compliance gate prevents hallucinated statistics that would land us in regulatory hot water (cf. RACQ ASIC action, Sep 2025).

*W Patch
start.*

Assumption 1 -- Causation language drives retention [CRITICAL]

- A/B test: split cohort -- control (current template) vs treatment (GenAI causation letter). MDE: 2pp churn reduction.
- Falsified if: no statistically significant difference ($p < 0.05$) after 8 weeks and 5,000+ notices per group.
- Falsification kills 85% of projected returns. Test this first.

Assumption 2 -- Data enables premium attribution [HIGH RISK]

- Target: $\geq 70\%$ of premium change attributable to tracked variables in ≥ 3 of 5 markets.
- Falsified if: legacy PAS cannot expose input-level rating variables for $\geq 50\%$ of policies.
- If fallback templates used for $> 50\%$ of policies, expected churn reduction drops to $\sim 0.5\text{pp}$ (-67% value).

Assumption 3 -- Regulatory acceptance [TIMELINE RISK]

- Submit sample notices to FCA Innovation Hub, BaFin AI-contact point, and equivalents in FR/IT/ES.
- Falsified if: ≥ 3 of 5 regulators mandate per-notice human review, or any enforcement action.
- Regulatory ambiguity alone could delay deployment 6-12 months -- 3 of 5 regulators have not clarified stance.

All three have clear falsification criteria. If any fails, the business case changes materially. Test order: 1 -> 2 -> 3, in parallel where possible.

7. Investment -- Cost Breakdown

Base scenario

Implementation (one-time)	EUR150K	McKinsey FS GenAI bench
Run cost (annual)	EUR20K	
Change management	EUR25K	Prosci change-mgmt bench
Model inference (415K renewals)	EUR7K	Anthropic pricing Jun 2026
Monitoring	EUR15K	

Base Year 1 Total: EUR217K

Range: Optimistic EUR113K | Base EUR217K | Pessimistic EUR390K

W Patch #2: Missing: Data readiness (Phase 0) -- EUR80K per market, 2-3 months. Not included in base budget.

W Patch #3: Missing: Regulatory delay costs (EUR50-100K legal, 12-18 months engagement timeline).

8. Returns -- Value Breakdown (Projected, Contingent)

W Patch #1: 85% of returns below depend on Assumption 1 (A/B test). These are projected values, not earned savings.

Annual projected value (if Assumptions 1-3 all hold)

Revenue retained (1.5pp churn red.)	EUR4.5M/yr	(85% of total)
Risk avoided (regulatory)	EUR750K/yr	(14%)
Call deflection	EUR34.5K/yr	(<1%)
Time saved	EUR51K/yr	(<1%)

Total projected: EUR5.34M annual (contingent on A/B validation)

Revenue retained is 85% of total value and depends entirely on Assumption 1 holding. 'Risk avoided' (EUR750K) has no published methodology -- it is an estimate of avoided ASIC-style enforcement cost, not a measured benchmark. 'Call deflection' and 'Time saved' are marginal (<1% each).

Base payback: 0.5 months

Pessimistic payback: 2.7 months

Top sensitivity: Churn reduction % and implementation cost

Phase deployment limits upfront risk. Even at pessimistic assumptions, payback is under one quarter. LLM variable cost (~EUR0.005/letter) is negligible vs retained premium.

W Patch #3: Regulatory-delay scenario added: EUR0 savings Year 1, EUR3.2M Year 2. Payback 14 months from first revenue. Scenario is missing from original analysis.

Sensitivity scenarios

Scenario	Payback	Annual Value	Driver
Base case	0.5 mo	EUR5.34M/yr	All assumptions hold
Pessimistic	2.7 mo	EUR3.80M/yr	Higher impl., 1.0pp churn red.
Regulatory delay*	14 mo*	EUR3.20M/yr	12-mo clearance delay, partial
Data failure	N/A	EUR0.75M/yr	Phase 1 fails >50% of policies

* Payback from first revenue, not project start. Realistic go-live: mid-2028.

Ask: EUR80K for a 4-week, two-market pilot (Phase 1a)

What the pilot tests:

- Assumption 1 -- does causation language drive retention? (A/B, 5,000+ per group)
- Assumption 2 -- can we decompose premium for $\geq 70\%$ of policies? (data audit)
- Regulatory sandbox submissions -- initial FCA and BaFin reactions (non-binding)

Go/no-go at Week 5:

- Positive: full rollout subject to Phase 0 (data readiness) and regulatory clearance funding
- Negative: EUR80K total sunk cost. Business case invalidated. Walk away.

W Patch #3: 6-month deployment target is from regulatory clearance, not project start. Including Phase 0 data readiness (2-3 mo) and regulatory engagement (12-18 mo est.), realistic go-live is mid-2028, not end-2026.

The pilot costs < EUR100K and answers the only question that kills the business case: do transparent notices actually retain customers in our specific book? If yes, we fund Phase 0 data readiness and regulatory engagement in parallel. If no, we walk away having spent less than one analyst-year.